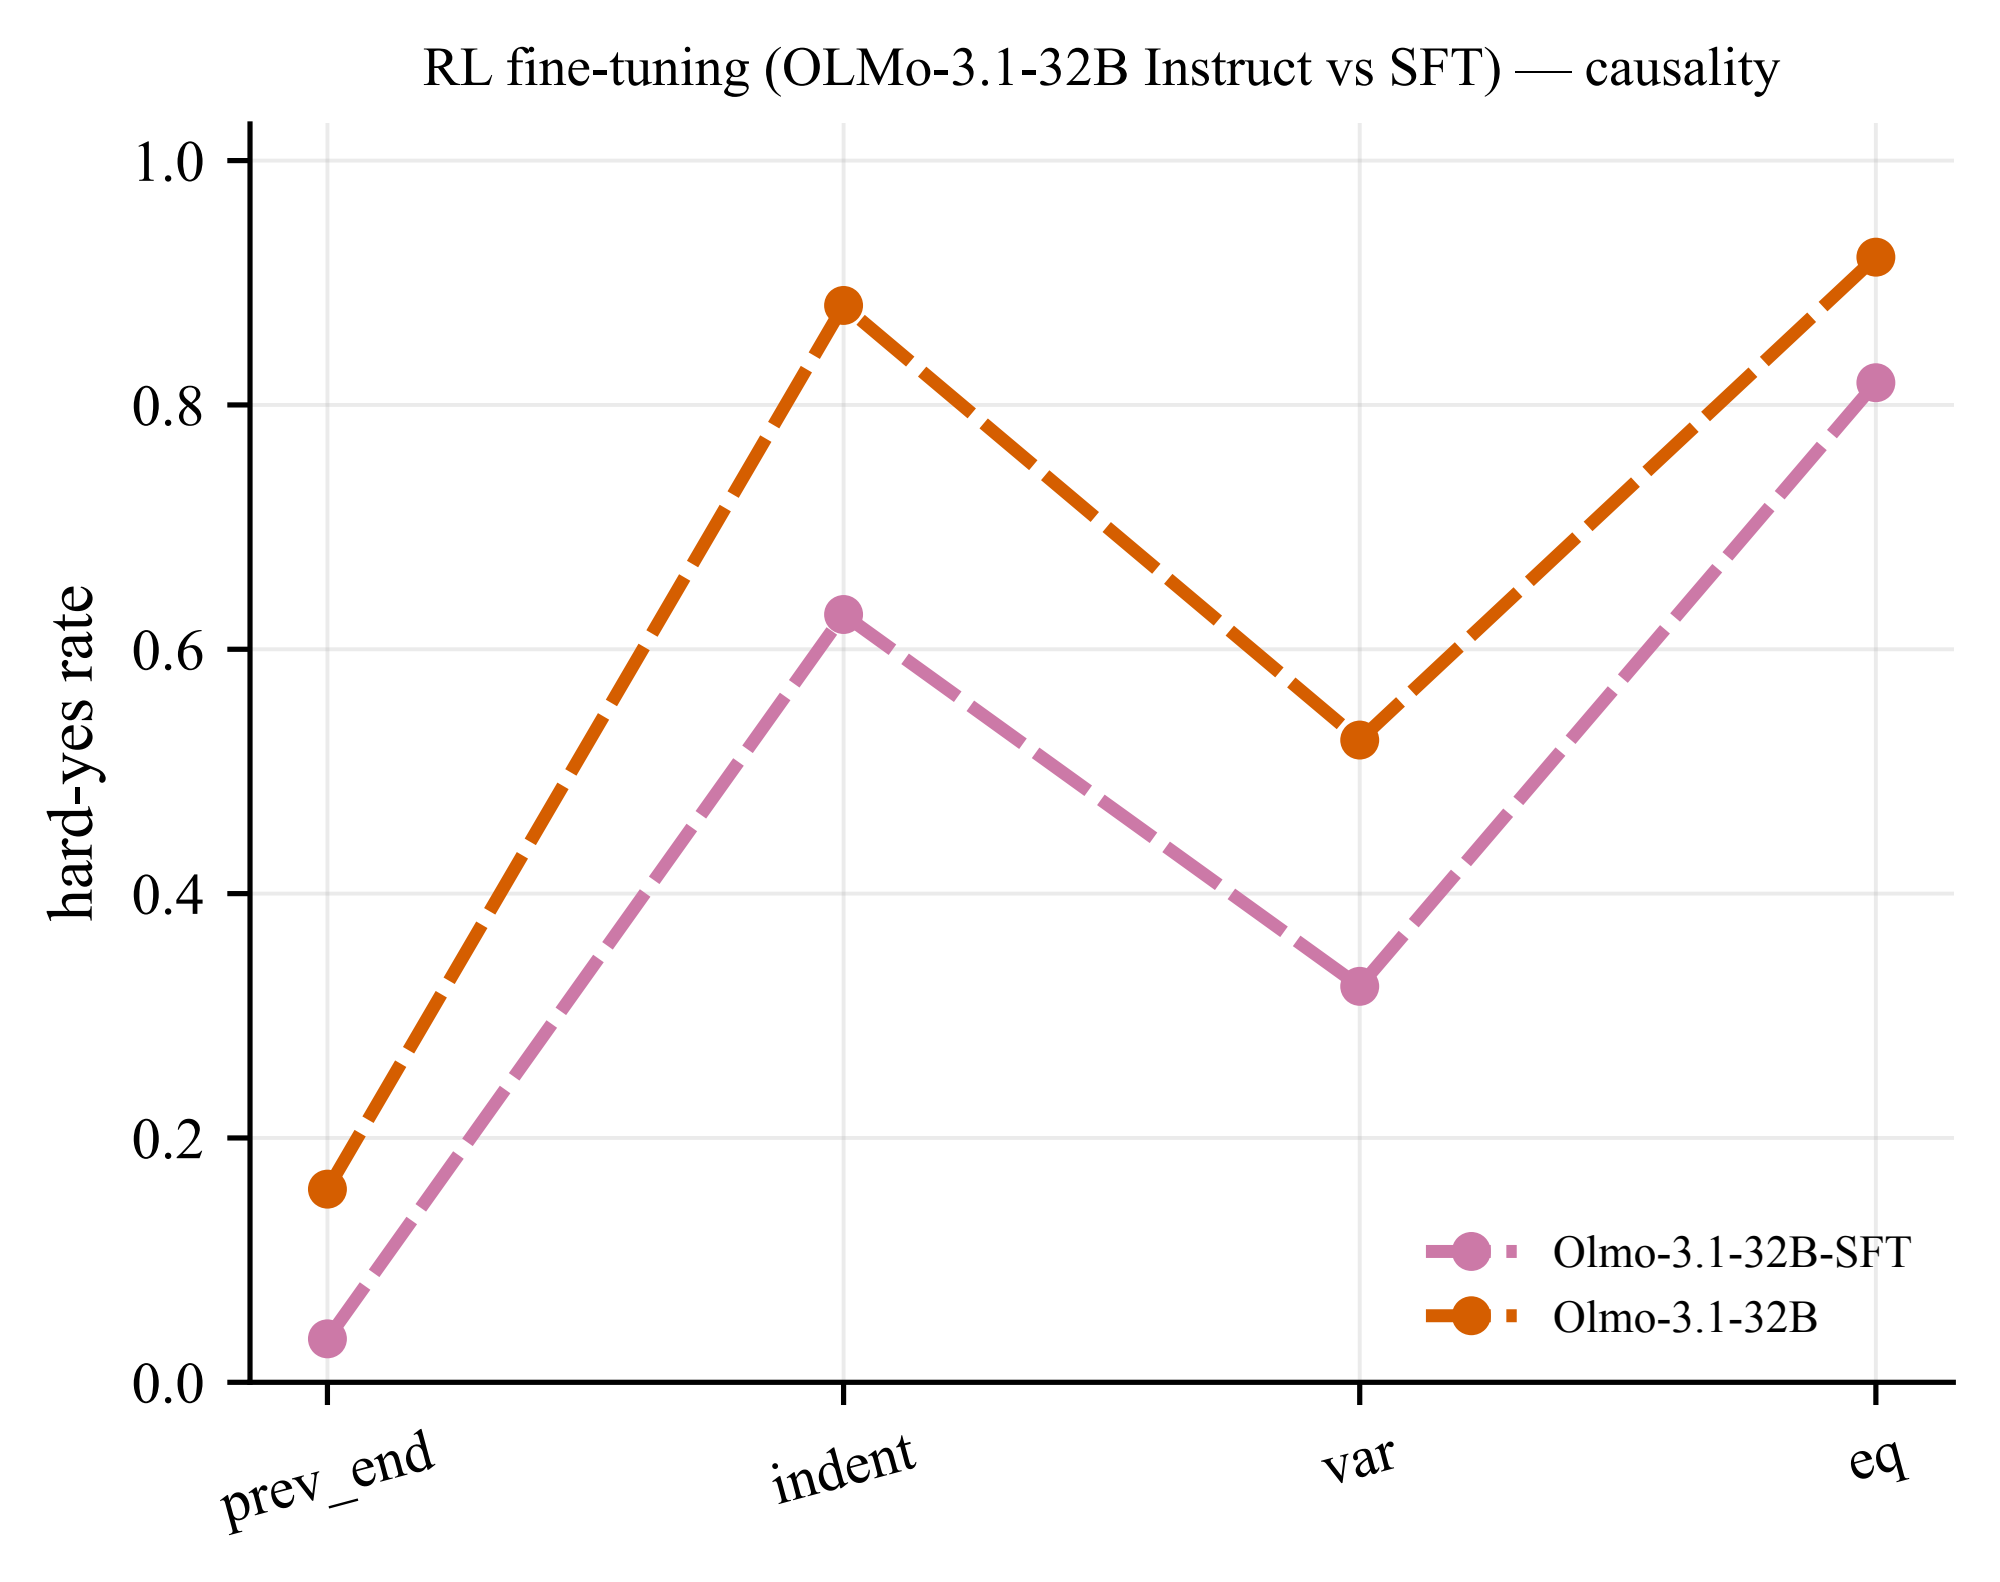
RL fine-tuning (OLMo-3.1-32B Instruct vs SFT) — causality
hard-yes rate
1.0
0.8
0.6
0.4
0.2
0.0
prev_end
indent
var
eq
Olmo-3.1-32B-SFT
Olmo-3.1-32B